

SIDDHANT KULKARNI

New York, USA | +1 347-646-9459 | sk10841@nyu.edu | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

EDUCATION

New York University, New York - Cumulative GPA: **3.68/4**

Aug 2023 - May 2025

Master's of Science in **Computer Science**

Savitribai Phule Pune University, India - Cumulative GPA: **9.38/10**

Jul 2019 - May 2023

Bachelor of Engineering in Information Technology (Honors **AI/ML**)

SKILLS

Languages: Python, C++, Java, JavaScript, SQL, **LLM & GenAI:** LangChain, LangGraph, OpenAI API, Azure AI Foundry, AWS Bedrock, Gemini, Llama, **Cloud:** AWS (Lambda, ECS, EKS, S3, DynamoDB, RDS), Redis, Postgres, **Frameworks:** FastAPI, Node.js, React, n8n, **Infra:** Docker, Kubernetes, Terraform, CI/CD.

EXPERIENCE

BEL Capital Advisory LLC (via Seawolf AI), NY - *AI Platform Engineer (Data and Cloud)*

Nov 2025 - Present

- Architected a production **multi-agent LLM pipeline** (retrieval, extraction, evaluation, KPI scoring agents) using **RAG and Gemini-2 Flash** on **AWS Lambda/ECS Fargate** to transform 1M+ SEC filings across 15 years into structured **governance intelligence** for institutional investors.
- Designed distributed **LLM inference pipelines** ingesting **EDGAR/FMP** feeds into **Postgres** (RDS) with asynchronous workers, supporting **40K+ document extractions/day** at **1.8s median latency** while reducing inference costs 30% through **caching and batching** optimizations.

Capital Group (via Seawolf AI), NY - *AI Engineer*

Jun 2025 - Nov 2025

- Developed a production **agent research assistant** using **GPT-4o** (Azure AI Foundry) that synthesized financial filings into investment theses, monitoring signals and automated research briefs used by **9,000+ analysts and 300+ PMs**.
- Achieved **92.4% precision / 89.7% recall** while engineering real-time financial signal pipelines using **AWS, Redis streams and DynamoDB**, reducing research turnaround from **hours to minutes** with 2s inference latency.

Jefferies Group LLC (via Seawolf AI), NY - *AI Engineer Intern*

Feb 2025 - May 2025

- Implemented an **LLM document intelligence** pipeline using **Claude 3.5 Sonnet** (AWS Bedrock) with Lambda, S3, SQS and DynamoDB to extract structured insights from financial documents.
- Operationalized the platform on Kubernetes (**EKS**) with Terraform and Bamboo CI/CD pipelines, scaling ingestion capacity to **10K+ documents per day** while achieving 87.8% extraction accuracy.

Siemens Digital Industries Software - *AI Engineer Intern*

Jun 2024 - Aug 2024

- Optimized **GPU-based LLM inference** pipelines for **CodeLlama 34B/70B**, improving latency and throughput for AI-assisted code understanding across large-scale C++ enterprise codebases.
- Performed **Static Code Analysis** for C++ with Coverity and created 15+ custom checkers in Codexm, boosting code maintainability by 20%.

Systematic Ventures, NY - *AI Engineer Intern*

Jun 2024 - Aug 2024

- Built a **RAG-based financial intelligence system** combining web scraping, embeddings and FAISS semantic search to retrieve insights across 3,500+ datasets for improving retrieval relevance for investment analysis workflows.
- Implemented a **vector search** index storing 1M+ embeddings, achieving **RAGAS** evaluation score of 0.72.

PROJECTS

TuneTether: Connect users through music vibes - *Apache Spark, OpenAI API, NodeJS, ReactJS, FastAPI*

- Constructed a recommendation engine using **Apache Spark ALS collaborative filtering** and FastAPI/React, generating personalized song recommendations from user listening data.

Jazzify: AI-driven Jazz Composition - *Tensorflow, Keras, RNN, Music21, LSTM*

- Trained an **LSTM generative model** on 78 jazz compositions to produce stylistically consistent music sequences with **94%** training accuracy.

ACTIVITIES

Teaching Assistant - Natural Language Processing, NYU Courant

Sep 2024 - Jan 2025

- Taught 150 students NLP and supervised 9 project teams implementing **HMMs**, **MT** systems, **POS/NER** pipelines, semantic role labeling using **NLTK** and **spaCy**.

Member - Vertically Integrated Projects, NYU

Sep 2024 - Dec 2025

- Developed time-series forecasting models (GPT-2, T5, Llama-2, LSTM) optimized for real-time inference efficiency on edge devices.